



LOAD BIG DATA EFFICIENTLY

PART 9: FASTER DATA LOADS IN A NUTSHELL



The File Format Matters

File format matters

Unstructured

text, e.g. txt

videos, e.g. mp4

sound, e.g. mp3

pictures, e.g. png

**Unorganised and
unformatted
information and data**

Semi-structured

{ j s o n }



xml

**Contains tags or markers
that create some structure
or semantic but not fully
enforced**

Structured



**Enforce schema and
data type rules,
tabular format**

File format matters

- CSV and JSON are both human readable, row level formats allowing fast writes and easy to open and process
- Compression is possible but not as fast as Avro and Parquet. It also limits the splitability
- CSV is further easy corruptable
- Avro is very efficient to save big amounts of data fast
- Parquet is very efficient in loading data including predicate pushdown but also writing data
- Compressable
- Avro and Parquet are not human readable and support is more limited especially for Avro

Size, write and loading times

10 Million row dataset with 8 files has been used

	Size	Monthly Costs
JSON	100 000 GB	2.000 €
AVRO	6 000 GB	120 €

Format	Size	Write time	Load time
JSON	1208 MB	3 s	6 s
CSV with infer schema	593 MB	3 s	12 s
PARQUET	81.5 MB	2 s	0,6 s
AVRO	69.2 MB	1 s	0,9 s

{j s o n}

vs.



- 94 % size

- 66 % write time

- 85 % load time



vs.



- 86 % size

- 33 % write time

- 95 % load time

Smaller size and faster reads and writes save you compute and storage costs

Detailed loading times

Format	Load time	of which meta data	of which actual load
JSON	6 s	3 s	3 s
CSV with infer schema	12 s	25 ms + 6 s	6 s
PARQUET	0.6 s	30 ms	0.6 s
AVRO	0.9 s	0 s	0.9 s

Small file Problem

Small file problem

10 Million row dataset with
Large: 8 files and Small: 100.000 files has been used

Format	Meta Time (S)	Loading Time (S)	Total (S)	Meta Time (L)	Loading Time (L)	Total (L)
JSON	21 s + 6,8 min	18 s	7,7 min	3 s	3 s	6 s
CSV	19 s + 86 ms + 6,9 min	30 s	7,7 min	25 ms + 6s	6 s	12 s
PARQUET	0,3 s	5,8 min	5,8 min	30 ms	0,6 s	0,6 s
AVRO	-	5,1 min	5,1min	-	0,9 s	0,9 s

Filter and write

col1	col2
1	A
2	B
3	A

col1	col2
4	B
5	B
6	B

col1	col2
7	A
8	A
9	B

col1	col2
10	B
11	A
12	A

Driver

- Assigns data to partitions
- Delegates the work to executors, meaning partitions to tasks
- Create execution plans
- Saves meta data about the files

Worker 1

Core1

Core2

Worker 2

Core1

Core2

Problem Cause of Small files

- File meta data stored in driver memory
- Efforts querying data including open file, closing file and checking the storage files and directories
- Scheduling overhead for delegating the partitions to tasks
- Reduced parallelism due to more created partitions
- Increased CPU costs due to serialisation and deserialisation

Solving the Problem

- Increase the file size. Either by checking the source options or having an intermediate process saving files in bigger files. I like Delta and you the bin packing (optimize) option
- Use file formats like Parquet or Avro with less meta data to be stored
- Reduce meta data by defining the schema
- Reduce number files per partition and thus increase number of partitions e.g. by using `maxPartitionBytes` or `openCostInBytes`

The force of the Schema

The force of the schema

```
ddl_schema = "id bigint, date date, timestamp timestamp, idstring string, idfirst  
string, idlast string"
```

```
spark_schema = t.StructType(  
    [  
        t.StructField('id', t.LongType(), True),  
        t.StructField('date', t.DateType(), True),  
        t.StructField('timestamp', t.TimestampType(), True),  
        t.StructField('idstring', t.StringType(), True),  
        t.StructField('idfirst', t.StringType(), True),  
        t.StructField('idlast', t.StringType(), True)  
    ]  
)
```

```
sdf_parquet = spark.read.format("parquet").schema(ddl_schema).load(path_parquet)
```

Results Absolute

Format	Experiment	Meta Time (S)	Loading Time (S)	Total (S)	Meta Time (L)	Loading Time (L)	Total (L)
JSON	w/o schema	21 s + 6,8 min	18 s	7,7 min	3 s	3 s	6 s
JSON	w schema	-	18 s	18 s	-	3 s	3 s
CSV	w/o schema	19 s + 86 ms + 6,9 min	30 s	7,7 min	25 ms + 6s	6 s	12 s
CSV	w schema	-	32 s	32 s	-	6 s	6 s
PARQUET	w/o schema	0,3 s	5,8 min	5,8 min	30 ms	0,6 s	0,6 s
PARQUET	w schema	-	33 s	33 s	-	0,4 s	0,4 s
AVRO	w/o schema	-	5,1 min	5,1min	-	0,9 s	0,9 s
AVRO	w schema	-	16 s	16 s	-	1 s	1 s

Results % to schema

Format	w/o vs with schema small	w/o vs with schema large	Small vs large w/o schema	Small vs large with schema
JSON	- 96 %	- 50 %	- 99 %	- 83 %
CSV	- 93 %	- 50 %	- 97 %	- 81 %
PARQUET	- 91 %	- 33 %	- 99,99 %	- 99 %
AVRO	- 95 %	- 0 %	- 99,99 %	- 94 %

Smaller files with Open Cost Per
Bytes

Open Cost Per Bytes to deal with small files

Open Cost Per Bytes

- Represents the cost of creating a new partition
- based on the config "spark.sql.files.openCostInBytes"
- defaults to 4 MB
- Technically it adds the cost, e.g. 4 MB, to each file which is called padding
- Official description: The estimated cost to open a file, measured by the number of bytes that could be scanned in the same time. This is used when putting multiple files into a partition. It is better to over-estimate, then the partitions with small files will be faster than partitions with bigger files (which is scheduled first). This configuration is effective only when using file-based sources such as Parquet, JSON and ORC.

Reference: <https://youtu.be/vkOxEHEKYhA>

Results of open Cost In Bytes

OpenCost MB	JSON	AVRO
1 MB	5,8 min	5,5 min
2 MB	21 s	40 s
4 MB	21 s	41 s
6 MB	21 s	42 s
8 MB	22 s	2,1 min
10 MB	1,8 min	3,7 min
w/o schema 4 MB	7,7 min	5,1 min

Reference: <https://youtu.be/C-ykod0Yyf4>

Predicate and Aggregate Pushdown


Predicate Pushdown

Name	Female	Age
Steven	False	33
Peter	False	47
Ana	True	21
Laura	True	37

Goal: Get all data WHERE Age > 35

```
sdf.filter(f.col("Age") > 35)
```

Predicate Pushdown

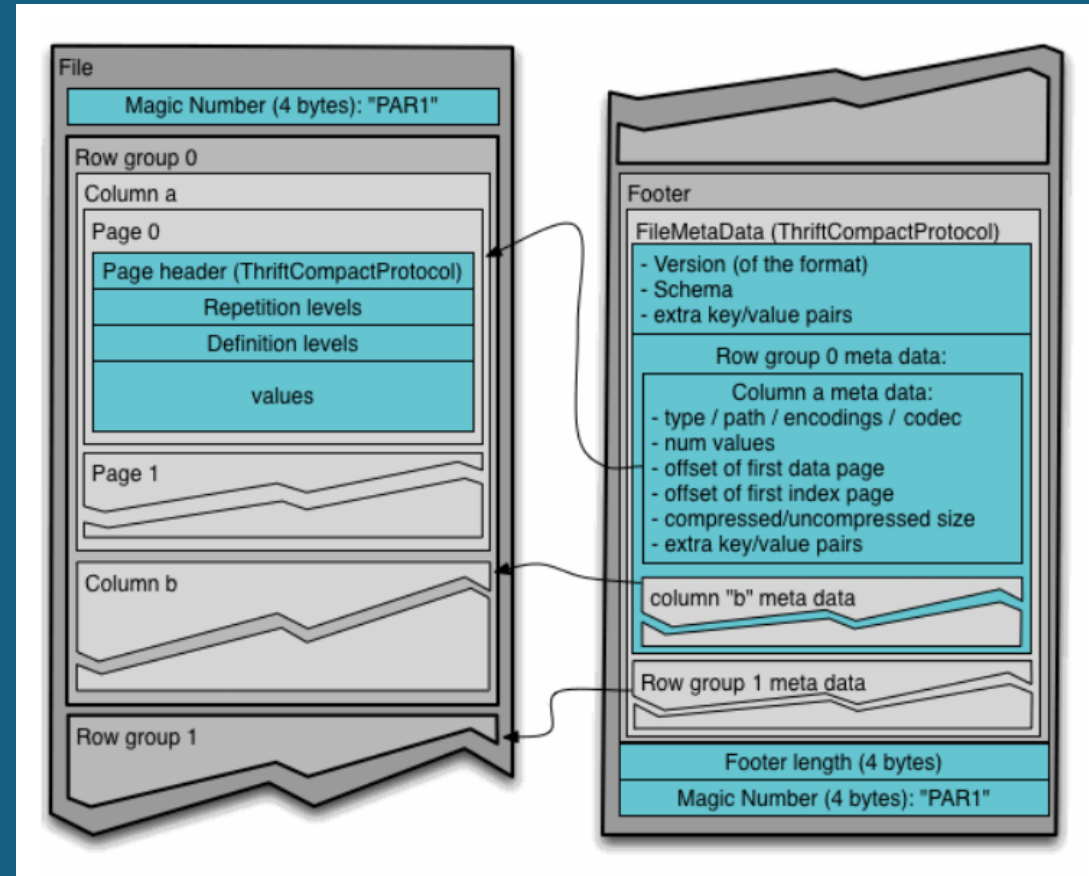


Predicate Pushdown

- **Predicate Pushdown is an optimization technique filtering data at the source and often relies on statistics**
- **Benefits:**
 - Less I/O meaning less data to load
 - Less memory usage
 - Faster queries
- **Parquet supports Predicate Pushdown using statistics saved in meta data footer**
- **Since Spark 3.1.0 also possible on Avro, CSV, JSON**



- Row Groups are a logical division on row level of a parquet defaulting to 128 MB
- Column part relates to column chunk of row groups
- Pages are invisible units where the encoding and compression happens
- Footer containing file metadata which can be used for predicate pushdown:
 - File level: num rows/ columns, schema
 - Row group: num rows/ columns
 - Column level: min, max, null count, distinct value counts, page indexes etc.



Load time and output rows for column/ row filters

Format	Load all data	Column filter	Row filter
JSON	16 s (10,000,000 rows)	5 s (10,000,000 rows)	4 s (300 rows)
CSV	13 s (10,000,000 rows))	4 s (10,000,000 rows)	4 s (308 rows)
PARQUET	2 s (10,000,000 rows)	0.6 s (10,000,000 rows)	0.1 s (20,000 rows)
AVRO	3 s (10,000,000 rows)	2 s (10,000,000 rows)	1 s (300 rows)

Aggregate Pushdown

- Filter and select push downs work on all data sources
- Aggregate Pushdowns work not on JSON, CSV, AVRO
- Activate aggregate pushdown for Parquet as follows:
 - `spark.conf.set("spark.sql.sources.useV1SourceList", "")`
 - `spark.conf.set("spark.sql.parquet.aggregatePushdown", "true")`
- Aggregate Pushdown has the following limitations:
 - No nested columns and string columns supported for min/max
 - Filter and aggregates are only for partitioned columns supported
- Aggregate Pushdown speeds up the performance significantly of counts, min and max
- V2 Source API seems unclear if more efficient than V1 but SQL interface seems different and Batch Scan is always on.

Better Partitions when loading data

Spark Architecture

col1	col2
1	A
2	B
3	A

col1	col2
4	B
5	B
6	B

col1	col2
7	A
8	A
9	B

col1	col2
10	B
11	A
12	A

Driver

- Assigns files to partitions
- Delegates the partitions as tasks to the worker
- Each core executes one task at the same time

Worker 1

Core1

Core2

Worker 2

Core1

Core2

What determines the number of partitions?

- Num of cores:
 - Spark tries to create at least the number of partitions equal to your number of cores
 - Can be changed with conf `spark.sql.files.minPartitionNum`
- Num of parquet files and its row groups: Parquet is only splittable on Row group level for partitioning
- Max Partition size:
 - Default 128 MB as the default row group size
 - Can be changed with conf `spark.sql.files.maxPartitionBytes`
- Max Cost per Bytes:
 - Represents the cost of creating a new partition, defaulting to 4 MB
 - Can be changed with conf `spark.sql.files.openCostInBytes`

Perfect distributions of Partitions



Bad example



Basic rules of good partitions

- Good parallelisation:
 - Factor 2-4 of your number of cores (exceptions for smaller files)
 - Uniform datasets generate also uniform partitions
- Partition size:
 - Too big partitions can lead to out of memory issues
 - Max partition size is at 128 MB, 100 MB to 1 GB is recommended
 - It depends of course on your machine and your other operations
- Distribution overhead:
 - A high number of partitions can create a distribution overhead
 - Execution time should make 90 % of the whole execution time
 - Exception: Small file problem where the distribution overhead is ok

Tip to check no. partitions: `sdf.rdd.getNumPartitions()`

Use Delta Lake



Key Features



ACID Transactions

Protect your data with serializability, the strongest level of isolation



Scalable Metadata

Handle petabyte-scale tables with billions of partitions and files with ease



Time Travel

Access/revert to earlier versions of data for audits, rollbacks, or reproduce



Open Source

Community driven, open standards, open protocol, open discussions



Unified Batch/Streaming

Exactly once semantics ingestion to backfill to interactive queries



Schema Evolution / Enforcement

Prevent bad data from causing data corruption



Audit History

Delta Lake log all change details providing a full audit trail



DML Operations

SQL, Scala/Java and Python APIs to merge, update and delete datasets

Summary

- Use Big data formats like Parquet and Avro
- Define the schema whenever possible
- Avoid small files if possible. If not reduce the overhead using the schema and conf `spark.sql.files.openCostInBytes`
- Apply column filter and row filter as close to the source as possible to use predicate pushdown
- Optionally the same for aggregate pushdown
- Use conf `spark.sql.files.maxPartitionBytes` to create a well distributed partitions